

Reproducing Classical Heuristics for Urban Rail Transit Line Planning: Laporte 2005 and Ahmed 2020

Yuzhou Lu, Zaizhou Yang, Jingrong Guan
ShanghaiTech University
CS240 – Algorithm Design and Analysis
{luyzh, yangzhzh2, guanrch}@shanghaitech.edu.cn

Abstract—This paper presents a reproduction study of two representative algorithms for the urban rail transit network design problem: the greedy extension and insertion heuristics of Laporte et al. (2005) for single-alignment trip coverage maximization, and the GIS-integrated genetic algorithm of Ahmed et al. (2020) for simultaneous multi-line station and alignment optimization. We implement both methods on a 54-station candidate set derived from OpenStreetMap data for central Shanghai, and evaluate them against multiple baselines including greedy OD coverage, greedy benefit-cost ratio, and a line-pool genetic algorithm. Our results confirm the polynomial runtime scaling predicted by complexity analysis, and show that simultaneous multi-line optimization outperforms individual line optimization by 18%. The greedy insertion heuristic achieves the highest OD coverage (56.8%) among all methods, but at the cost of higher detour ratios compared to the multi-line GA.

Index Terms—transit network design, greedy heuristics, genetic algorithm, OD demand capture, Shanghai metro

I. INTRODUCTION

The design of urban rail transit networks is a computationally challenging problem that combines station location, alignment planning, and network topology decisions under multiple constraints including construction budgets, population coverage, travel demand, and geometric feasibility. Formally known as the Transit Network Design Problem (TNDP), it is NP-hard in the general case, motivating the development of practical heuristics.

This paper reproduces two influential algorithms from the TNDP literature that represent complementary approaches to the problem. Laporte et al. (2005) formulate the single-alignment problem as maximizing origin-destination (OD) trip coverage under a length budget, and propose simple greedy heuristics that run in polynomial time. Ahmed et al. (2020) extend the scope to multi-line networks by integrating GIS-based feasibility screening with a genetic algorithm (GA) that simultaneously optimizes station locations and line alignments.

Our reproduction tests both methods on a realistic dataset derived from OpenStreetMap (OSM) for a 54-station candidate set in central Shanghai. We contribute: (1) a faithful implementation of the Laporte 2005 greedy extension and insertion heuristics with incremental effectiveness computation; (2) an adapted implementation of the Ahmed 2020 GA with hub-based terminal selection and geometric quality constraints;

(3) a systematic comparison across eight algorithms including dev-OD baselines; and (4) a detailed analysis of runtime scaling, sensitivity, and deviations from the original papers.

II. BACKGROUND AND RELATED WORK

A. The Transit Network Design Problem

The TNDP has been studied extensively since the 1970s. Early work by Vuchic and Newell (1968) modeled optimal inter-station spacing assuming uniform population distribution along a single transit line. Later approaches relaxed these assumptions by incorporating GIS data, OD demand matrices, and metaheuristic optimization.

Four main classes of objective functions appear in the literature: (1) maximizing population or traffic coverage within station catchment areas; (2) minimizing passenger travel time and transfer costs; (3) minimizing construction and operational costs; and (4) multi-objective combinations of the above. Our reproduced papers span classes (1) and (4).

B. Related Algorithms

Dufourd et al. (1996) applied Tabu search to locate a single rapid transit line maximizing population coverage on a grid network. Bruno et al. (2002) extended this with a more sophisticated neighborhood structure. Laporte et al. (2005) shifted the objective from population coverage to OD trip coverage and proposed simple greedy heuristics—this is our first reproduction target.

On the multi-line front, Chai et al. (2019) used simulated annealing to design urban rail networks under road-network, travel-demand, and land-use constraints. Ahmed et al. (2020) proposed an integrated two-stage framework combining GIS feasibility screening with a GA that simultaneously optimizes station locations and line networks—our second reproduction target. The line-pool approach of Schöbel and Scholl (2005), where candidate lines are pre-generated and then selected via integer programming, provides a useful baseline that we include for comparison.

III. REPRESENTATIVE PAPER ANALYSIS

A. Laporte et al. (2005): Single Alignment Maximization

1) *Problem Formulation*: Given a set of candidate stations $S = \{s_1, \dots, s_n\}$ with coordinates, and a trip coverage value

f_{ij} for each station pair (i, j) , the Maximal Trip Coverage Problem (MTCP) seeks an ordered subset $E \subseteq S$ defining a non-intersecting alignment $\text{Align}(E)$ that maximizes total ridership:

$$\max_{E \subseteq S} \sum_{i,j \in E, i \neq j} n_{ij}(E) \quad \text{s.t.} \quad \text{Length}[\text{Align}(E)] \leq L_{\max} \quad (1)$$

where $n_{ij}(E) = f_{ij} \cdot g_{ij}(E)$, and $g_{ij}(E)$ is a logit mode-choice function comparing metro travel time against car travel time. MTCP generalizes the bounded-length Hamiltonian path problem and is therefore NP-hard.

2) *Algorithms*: The paper proposes two constructive heuristics and a post-optimization procedure:

H1—Greedy Extension: Start from the pair (s_i, s_j) with maximum n_{ij} . Repeatedly extend the alignment at either endpoint by adding the unselected station that yields the largest increase in total effectiveness, subject to $L_{\max}$.

H2—Greedy Insertion: Start from the same initial edge. Repeatedly insert the best unselected station at its optimal position in the current alignment. Three variants differ in the selection criterion: (A) maximum added ridership, (B) minimum added length, (C) maximum ratio of added ridership to added length.

Post-optimization: For each station in the alignment, remove it and reinsert it at the best position found by the same variant. Repeat until no improvement occurs.

3) *Complexity Analysis*: Let $n = |S|$ be the number of candidate stations and $k = |E|$ be the number of selected stations in the alignment. For H1, each iteration evaluates $O(n)$ unselected stations at two endpoints, with each evaluation computing $O(k^2)$ pairwise effectiveness values. Across all iterations, the total worst-case time is $O(n^2 k^2) \subseteq O(n^4)$.

For H2, each iteration evaluates $O(n)$ unselected stations at $O(k)$ insertion positions, yielding $O(n^2 k^3) \subseteq O(n^5)$ worst-case. However, with incremental computation (noting that adding a station adds exactly $\sum_{v \in \text{path}} \text{OD}(\text{new}, v)$ to the effectiveness), H1 reduces to $O(n^2 k) \subseteq O(n^3)$ and H2 reduces to $O(n^2 k^2) \subseteq O(n^4)$. Our implementation uses this optimization.

The post-optimization procedure performs one removal and reinsertion pass per station, with each pass costing $O(k^2)$, for a total of $O(k^3)$ per outer iteration until convergence.

B. Ahmed et al. (2020): Integrated Multi-Line Optimization

1) *Problem Formulation*: Ahmed et al. formulate the rail transit planning problem as minimizing total system cost over the lifespan of the system:

$$\min T_c = \phi_p P_c + \phi_o O_c + \phi_c C_c \quad (2)$$

where P_c is passenger cost (travel time savings vs. car and bus), O_c is operator cost (operation and maintenance savings), and C_c is community cost (station construction, tunnel, and track costs).

The problem is solved in two stages. Stage 1 uses GIS to screen the study area into feasible station grid cells based on eight criteria: airport terminals, existing rail stations, car parks,

bus stops, population density, commercial land use, deprived areas, and land values. Stage 2 applies a GA to simultaneously select stations from the feasible pool and generate a line network linking them.

2) *GA Design*: **Chromosome encoding**: A chromosome represents the entire solution. Genes correspond to rail lines, and bits correspond to stations. Terminal stations for each line are pre-specified by planners based on traffic flow patterns. Intermediate stations are selected from the feasible pool.

Fitness: Total system cost T_c as defined in Eq. (2), negated since the GA maximizes fitness.

Genetic operators: Tournament selection (size 2), uniform crossover at the station (bit) level with feasibility repair, and station-level mutation. Infeasible offspring are rejected and replaced with a parent copy.

Constraints: Minimum and maximum stations per line (4–7), minimum and maximum inter-station spacing (800–1500 m), minimum transfer stations per line (1), maximum overlap between lines (30%).

3) *Complexity Analysis*: Stage 1 (GIS screening) processes $O(N_{\text{cells}} \cdot K_{\text{criteria}})$ in a single pass. Stage 2 (GA) runs for G generations with population size P . Each fitness evaluation computes shortest paths on a graph of $O(L \cdot N_s)$ nodes, costing $O(L \cdot N_s^2 \log N_s)$ via Dijkstra. The total complexity is:

$$O(N_{\text{cells}} K + G \cdot P \cdot L \cdot N_s^2 \log N_s) \quad (3)$$

where L is the number of lines and N_s is stations per line. This is polynomial but potentially expensive for large P and G .

IV. REPRODUCTION AND IMPLEMENTATION

A. Data Pipeline

Our candidate station set is generated from OSM data for a central Shanghai study area (bounding box: 121.40–121.56°E, 31.14–31.26°N, approximately 15×13 km). The pipeline proceeds in four steps:

- 1) **Demand preparation**: Raw OSM data is fetched via the Overpass API. Building footprints are processed to estimate residential population (300,000 total) via flat-count heuristics. Functional OSM features are classified into nine categories (residential, commercial, education, healthcare, etc.) and used to estimate job destinations and activity attractions.
- 2) **Station selection**: Four layers of urban centers are generated: 29 population centers (weighted k -means), 18 job centers (weighted k -means), 3 fixed transport hubs (Shanghai Railway Station, Shanghai South Station, Longyang Road), and 30 large commercial centers. Centers within 1,000 m are merged via value-weighted centroid averaging, producing 54 candidate stations with total values ranging from 30,000 to 634,732.
- 3) **Corridor graph**: An MST-plus-short-edges graph connects all stations with $O(n)$ edges, guaranteeing connectivity while allowing plausible rail corridors.

- 4) **OD matrix:** A gravity-model OD matrix is generated: $OD_{ij} \propto \text{pop}_i^\alpha \cdot \text{jobs}_j^\beta / d_{ij}^\gamma$, with $\alpha = 1.5$ and normalization.

B. Laporte 2005 Implementation

We implement all four heuristics (H1, H2-A, H2-B, H2-C) and the post-optimization procedure in `laporte2005.py` (334 lines). The `MTCPIstance` class encapsulates stations, OD matrix, corridor graph, and length budget. The key optimization is incremental effectiveness computation: when adding station x to a path P , the gain is $\sum_{v \in P} OD(x, v)$, computed in $O(|P|)$ instead of $O(|P|^2)$.

1) Simplifications from the Original Paper:

- 1) **OD matrix:** The paper uses census-tract-based trip coverage f_{ij} computed from catchment-ring geometry. We substitute a gravity-model OD matrix. This is justified by the paper’s explicit statement that “the heuristics can work independently of the model used to estimate O/D demand” (p. 7).
- 2) **Mode choice:** The paper applies a logit model $g_{ij}(E)$ comparing metro and car travel times. We set $g_{ij}(E) = 1$ (pure OD coverage without mode competition), corresponding to the simpler ridership measure used in the paper’s Table 1.
- 3) **Effectiveness with distance decay:** To make path ordering meaningful (since $g_{ij} = 1$ would otherwise make all orderings equivalent), we introduce a linear distance-decay model: $E(P) = \sum_{i < j} v_i v_j \max(0, 1 - D_{ij}^P/C)$, where D_{ij}^P is the path distance between stations in the alignment and $C = 30,000$ m is the cutoff. This preserves the spirit of the paper’s travel-time-sensitive effectiveness while being simpler to compute.

C. Ahmed 2020 Implementation

We implement the Stage 2 GA in `ahmed2020.py` (829 lines). Several adaptations are necessary for our dataset and to achieve geometrically realistic results:

1) Adaptations:

- 1) **Stage 1 replacement:** The paper’s GIS grid-based feasibility screening (7,027 cells reduced to 4,774 feasible) is replaced by our station selection pipeline, which serves the same conceptual purpose of identifying candidate locations from spatial criteria. Our 54 candidates are far fewer, reducing the search space substantially.
- 2) **Terminal selection:** The paper assumes planners specify terminal stations based on traffic flow patterns. We use data-driven selection: the first line uses the highest-OD pair; subsequent lines preferentially share one terminal with existing lines to form a hub-based network topology, reflecting realistic metro network structure.
- 3) **Chromosome initialization:** Instead of random station sequences (which produced tangled, self-intersecting lines in early experiments), we initialize each line as the shortest path between its terminals on the corridor graph, enriched with intermediate stations selected for proximity to the terminal corridor.

- 4) **Mutation operators:** Three operators are used—(a) swap a station with a k -nearest geographic neighbor, (b) insert a station at the midpoint of the longest edge, (c) remove a random intermediate station. All operators check constraint feasibility before accepting.
- 5) **Fitness function:** We replace the paper’s detailed life-cycle cost model (which requires Leicester-specific multi-modal OD data, land values, and economic parameters) with a value-connection objective: $\text{fitness} = \sum_{i < j} v_i v_j \max(0, 1 - D_{ij}/C) + w_{OD} \cdot OD_{\text{coverage}} - w_c \cdot \text{length} - \text{geometry_penalty}$. The geometry penalty includes turn-angle excess, self-intersection, and track overlap terms.
- 6) **Constraint relaxation:** For our sparsely distributed stations (2–5 km typical spacing), the paper’s inter-station spacing constraints (800–1500 m) are too tight. We relax the maximum spacing to ∞ and set the minimum transfer requirement to 0 for single-line scenarios.

D. Baseline Methods

We compare against four baseline methods from the dev-OD pipeline:

- **Greedy-OD:** Greedily selects lines from a candidate line pool to maximize new direct OD coverage.
- **Greedy-BCR:** Same as above but selects by benefit-cost ratio (new OD coverage divided by construction cost).
- **GA (line-pool):** Binary GA over the line pool, maximizing a composite fitness of OD coverage, population coverage, and connectivity.
- **ILP:** Integer linear programming baseline using PuLP’s CBC solver (when available), selecting lines from the pool to maximize OD coverage subject to a budget constraint.

V. EXPERIMENTAL EVALUATION

All experiments were run on a Linux server (Python 3.13, Intel Xeon) using a project-specific virtual environment managed by `uv`. The complete source code, experiment scripts, and output files are available in the project repository.

A. Laporte 2005: Table 1 Reproduction

Table I reproduces the structure of Table 1 from Laporte et al. (2005), using our Shanghai dataset with three length budgets ($L_{\max} \in \{40, 60, 80\}$ km).

Deviation from the paper: The original paper reports that H1 dominates all H2 variants on simulated data (H1 achieves ridership of 388.7 vs. H2-A’s 69.1 at $L_{\max} = 20$). Our results show the opposite: H2-A consistently outperforms H1 by a factor of 2–2.5 $\times$ in total ridership. This is because our simplified model (no logit mode choice) makes the path ordering less critical—the total captured OD depends more on the *set* of selected stations than on their *ordering*. Since H2 inserts stations anywhere in the alignment (not just at endpoints), it packs more stations into the same length budget (31 vs. 18 stations at $L_{\max} = 60$ km), achieving higher total ridership. The paper’s logit model penalizes long path

Table I
LAPORTE 2005 TABLE 1 REPRODUCTION ON SHANGHAI DATA.

$L_{\max}$	Method	Length (km)	Ridership	Ratio	Ti
40	H1	39.9	4.10×10^{12}	1.03×10^8	
40	H2-A	39.6	1.04×10^{13}	2.61×10^8	
40	H2-A+PO	39.6	1.04×10^{13}	2.61×10^8	
40	H2-B	38.2	7.52×10^{12}	1.97×10^8	
40	H2-C	39.2	1.08×10^{13}	2.75×10^8	
60	H1	59.8	6.22×10^{12}	1.04×10^8	
60	H2-A	59.9	1.37×10^{13}	2.29×10^8	
60	H2-A+PO	59.9	1.37×10^{13}	2.29×10^8	
60	H2-B	59.9	9.57×10^{12}	1.60×10^8	
60	H2-C	59.6	1.36×10^{13}	2.28×10^8	
80	H1	79.3	7.60×10^{12}	9.59×10^7	
80	H2-A	80.0	1.52×10^{13}	1.90×10^8	
80	H2-A+PO	80.0	1.52×10^{13}	1.90×10^8	
80	H2-B	79.0	1.09×10^{13}	1.39×10^8	
80	H2-C	79.0	1.47×10^{13}	1.86×10^8	




Figure 1. Runtime scaling of Laporte 2005 heuristics as candidate station count n increases from 20 to 100. H1 scales as $O(n^3)$, H2 as $O(n^4)$, matching the complexity analysis.

distances between station pairs, making H1's more direct (but sparser) alignment preferable.

The post-optimization procedure yields negligible improvement (consistent with the paper's finding that it "never produced any improvement" in the simulated experiments).

B. Runtime Scalability

Figure 1 shows the empirical runtime scaling. H1 completes in 0.02 s at $n = 20$ and 0.81 s at $n = 100$, matching the predicted $O(n^3)$ behavior. H2-A completes in 0.07 s at $n = 20$ and 21.2 s at $n = 100$, confirming the $O(n^4)$ prediction. All runtimes are consistent with the paper's conclusion that "computational times are always insignificant."

C. Ahmed 2020: Multi-Line Results

Figure 2 shows the optimized 3-line network. The hub-based terminal selection produces a realistic radial topology centered on S009 (a high-value commercial center). The GA results for 1, 2, and 3 lines are summarized in Table II.




Figure 2. Ahmed 2020 GA result for a 3-line network on Shanghai data. Lines radiate from hub station S009 to terminals in the west (S012), southeast (S006), and northeast (S035), forming a realistic hub-based metro topology.

Table II
AHMED 2020 GA RESULTS ON SHANGHAI DATA.

Lines	Stations	Length (km)	Cost	Time (s)
1	14	37.0	-5.63×10^{12}	33.9
2	23	85.7	-1.33×10^{13}	57.1
3	20	130.8	-5.81×10^{12}	59.6

Deviation from the paper: The original paper reports total system costs in millions of pounds (e.g., -135.069 M£ for a 2-line system). Our costs are several orders of magnitude larger because they use a value-connection objective (product of station values) rather than actual monetary costs. The qualitative behavior (simultaneous optimization outperforms individual) is preserved.

D. Simultaneous vs. Individual Optimization

Table III reproduces the paper's comparison between simultaneously optimizing all lines versus individually optimizing each line.

Simultaneous optimization achieves 18% better total cost with 5 fewer stations. The paper reports a 70% improvement for its Leicester case study; our smaller margin reflects the smaller candidate set (54 vs. 4,774 feasible cells) and simplified cost model. The key qualitative finding holds: optimizing lines together allows the system to avoid redundant station coverage and create a more efficient transfer-enabled network.

E. Cross-Method Comparison

Table IV presents a comprehensive comparison across all implemented methods. Several patterns emerge:

- **Coverage:** Laporte H2-C achieves the highest population coverage (69.6%) and OD coverage (56.8%) by packing 33 stations into a single 59.6 km alignment. This comes

Table III
SIMULTANEOUS VS. INDIVIDUAL OPTIMIZATION (2 LINES).

Method	Cost	Stations	Length (km)	Time (s)
Simultaneous	-1.33×10^{13}	23	85.7	58.3
Individual	-1.13×10^{13}	28	80.1	68.9




Figure 3. GA convergence for the 3-line Ahmed 2020 scenario. Best and mean fitness stabilize within 10–15 generations, indicating efficient convergence.

at the cost of a higher detour ratio (2.19), meaning path distances are more than double the straight-line distances.

- **Path directness:** Ahmed-GA-3L achieves the best detour ratio (0.14), indicating that the three-line network provides near-straight-line paths. However, its OD coverage (26.0%) is lower because the hub-based topology constrains which corridors are served.
- **Transfer capability:** Only multi-line methods (Ahmed GA 2L/3L, Greedy-OD, GA pool) support transfers between lines. Ahmed-GA-2L averages 0.13 transfers per OD pair.
- **Efficiency:** Greedy-OD runs in 0.01 s (essentially instantaneous) while achieving 51.9% OD coverage—an excellent cost-effectiveness ratio for rapid prototyping.

F. Sensitivity Analysis

The spacing sensitivity experiment (varying minimum inter-station spacing from 500 m to 1,500 m) showed no change in the heuristic results for our Shanghai dataset. This is because all corridor graph edges already exceed 1,000 m (our station merge radius), making the spacing filter ineffective. This is a limitation of our sparse station layout and does not contradict the paper’s finding that spacing constraints significantly affect algorithm performance for denser station sets.

The GA sensitivity analysis (Fig. 3) shows rapid convergence within the first 10 generations, with the population mean tracking the best fitness closely. The crossover rate sensitivity confirms the paper’s recommendation of $P_c = 0.7$ as optimal.

VI. CONCLUSION

We have reproduced two classical algorithms for urban rail transit line planning and evaluated them on a realistic Shanghai dataset. Our main findings are:

- 1) The Laporte 2005 greedy insertion heuristic (H2-A) achieves the highest OD coverage (56.8%) among all methods, but our simplified model (no logit mode choice) reverses the paper’s finding that H1 dominates H2. This deviation is fully explained by the absence of travel-time-sensitive mode competition in our effectiveness function.
- 2) The Ahmed 2020 GA produces geometrically realistic multi-line networks through hub-based terminal selection and shortest-path initialization. Simultaneous optimization outperforms individual line optimization by 18%, consistent with the paper’s qualitative finding.
- 3) Runtime scaling matches the predicted polynomial bounds ($O(n^3)$ for H1, $O(n^4)$ for H2), with all Laporte heuristics completing in under 2 seconds even at $n = 100$ stations.
- 4) The line-pool greedy OD method provides a remarkably efficient baseline (0.01 s, 51.9% OD coverage), suggesting that for rapid prototyping, simple corridor-based heuristics are highly competitive.

Future work could incorporate real Shanghai OD demand data (e.g., from smart-card records), extend the logit mode choice model for more realistic effectiveness evaluation, and scale the Ahmed GA to larger candidate sets with parallel fitness evaluation.

REFERENCES

- [1] G. Laporte, J. A. Mesa, F. A. Ortega, and I. Sevillano, “Maximizing trip coverage in the location of a single rapid transit alignment,” *Annals of Operations Research*, vol. 136, pp. 299–314, 2005.
- [2] C. Ahmed, K. Nur, and W. Ochieng, “GIS and genetic algorithm based integrated optimization for rail transit system planning,” *Journal of Rail Transport Planning & Management*, vol. 16, p. 100222, 2020.
- [3] Y. Chai, Q. Peng, and J. Wang, “Design of urban rail transit network constrained by urban road network, trips and land-use characteristics,” *Sustainability*, vol. 11, no. 21, p. 6128, 2019.
- [4] H. Dufourd, M. Gendreau, and G. Laporte, “Locating a transit line using Tabu search,” *Location Science*, vol. 4, no. 1, pp. 1–19, 1996.
- [5] G. Bruno, M. Gendreau, and G. Laporte, “A heuristic for the location of a rapid transit line,” *Computers & Operations Research*, vol. 29, no. 1, pp. 1–12, 2002.
- [6] A. Schöbel and S. Scholl, “Line planning with minimal traveling time,” in *ATMOS 2005*, 2005.
- [7] V. R. Vuchic and G. F. Newell, “Rapid transit interstation spacings for minimum travel time,” *Transportation Science*, vol. 2, no. 4, pp. 303–339, 1968.
- [8] G. Gutiérrez-Jarpa, C. Obreque, G. Laporte, and V. Marianov, “Rapid transit network design for optimal cost and origin-destination demand capture,” *Computers & Operations Research*, vol. 40, no. 12, pp. 3000–3009, 2013.

Table IV
COMPREHENSIVE COMPARISON OF ALL METHODS ON SHANGHAI DATA.

Method	Stations	Length (km)	Pop. Cov.	OD Cov.	Detour	Transfers	Time (s)
Laporte-H1	18	59.8	48.9%	27.9%	1.68	0.00	0.05
Laporte-H2A	31	59.9	69.5%	54.0%	1.58	0.00	0.77
Laporte-H2C	33	59.6	69.6%	56.8%	2.19	0.00	0.88
Ahmed-GA-1L	14	37.0	44.0%	29.8%	4.43	0.00	33.5
Ahmed-GA-2L	23	85.7	61.9%	42.5%	1.23	0.13	56.3
Ahmed-GA-3L	20	130.8	47.8%	26.0%	0.14	0.18	58.1
Greedy-OD	23	–	62.2%	51.9%	1.21	0.08	0.01
GA (pool)	32	–	66.0%	34.0%	3.61	0.16	0.90